

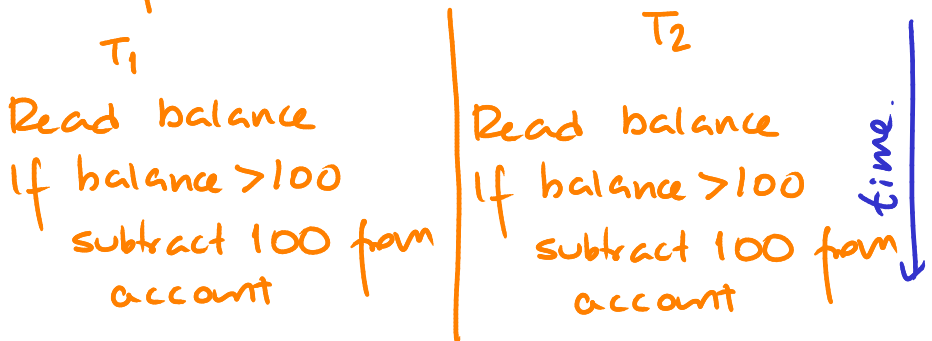
# Transactions and Concurrency

In production DBs users perform transactions concurrently.

- DBMS wants to maximize throughput
- Without compromising integrity

Example:

Person tries to remove, at the same time \$100 from bank account.



Can we have reach a state where person gets \$200 but bank only records \$100 given?

If so, we have lost consistency of data.

# Properties of Transactions:

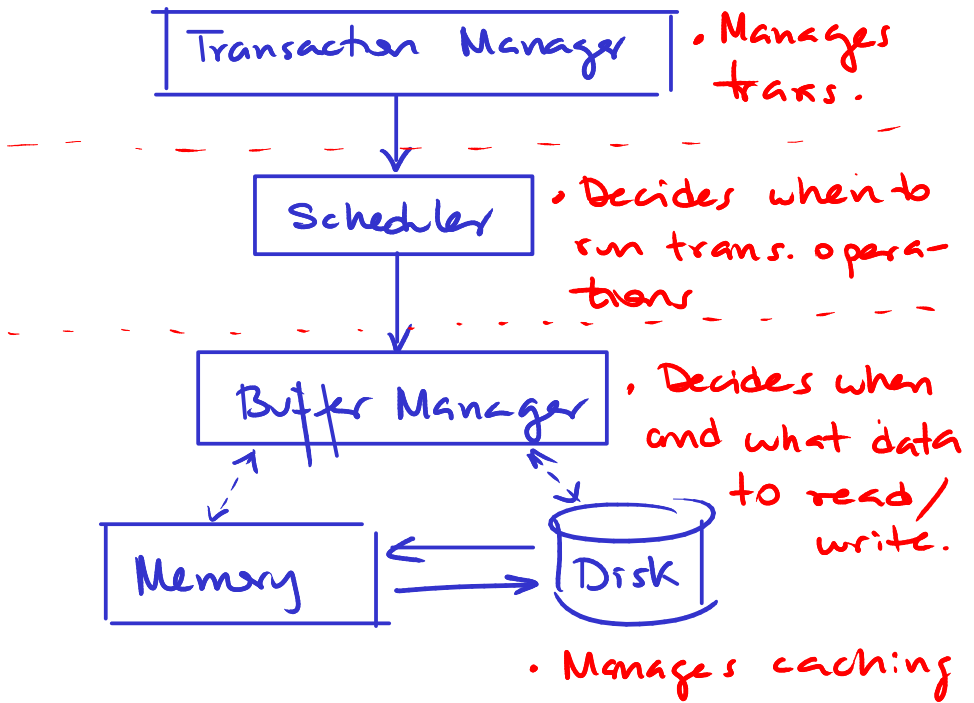
## ACID

- Atomicity: A transaction happens in its entirety or not at all.  
Incomplete transactions must be undone.
- Isolation: A transaction must appear to be executed as if no other transaction is executing at the same time.  
Transactions cannot communicate with each other.
- Durability: The effect on the db of a transaction has successfully completed must never be lost.  
• Even in the event of failures.

## Responsibility of Programmer.

- Consistency: Transactions are given a DB in a consistent state and are expected to keep it consistent.

The role of the DBMS is to maximize number of concurrent trans. while maintaining ACID.



To maximize throughput, the scheduler might:

- delay transaction operations.
- reorder transactions

To guarantee ACID, the scheduler:

- must make sure transactions are durable
- avoid undesirable interleaving of trans.
- deal with deadlocks

## Transactions

Any transaction either

**Atomicity.**

- completes (commits)
- or
- aborts (rollback)

If system crashes: (server or client):

- Non completed transactions must be undone (rollback) **Durability.**

## Correctness Principle

Any transaction, if executed in isolation will transform any consistent state of the DB into another consistent state.

The DBMS must guarantee isolation even when many trans. are executed concurrently

A transaction is a list of actions.  
For simplicity sake we will only consider read/write of DB objects.

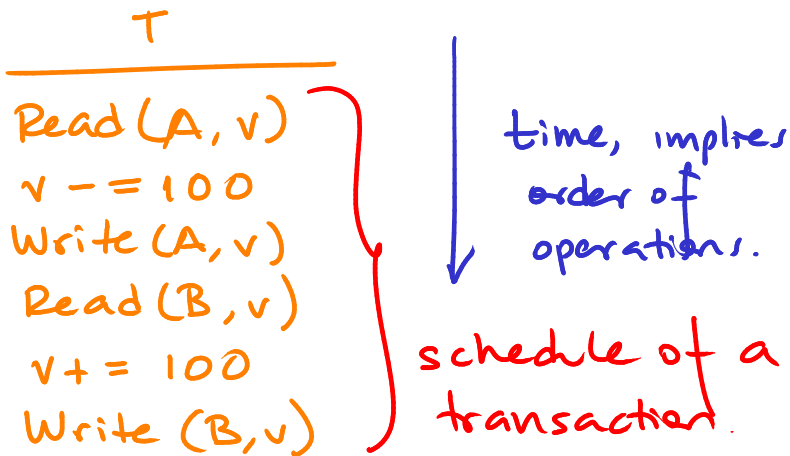
## Notation.

Read ( $A, v$ )

Reads DB object  $A$   
into local variable  $v$   
(local to transaction)

Write (A, v) Replaces DB object A with value in v.

Ex: T is a transaction that moves \$100 from account A to account B:



There might be many copies of the same transaction running.

Ex: Two instances of T are trying to run simultaneously.

Assumption:

Reads and writes are atomic and cannot be interleaved.

Schedule

Sequence of actions taken by one or more transactions.

When two transactions want to be executed 3 options:

1)  $T_1$  executes first, then  $T_2$   
denoted:

2)  $T_1; T_2$   
 $T_2; T_1$  } Serial schedules.

3) The operations of  $T_1$  and  $T_2$  interleave.

Many, many possible interleavings of operations of  $T_1$  and  $T_2$

- Some safe
- Some unsafe (break consistency)

### Serial Schedule

A schedule is serial if its actions consists of all the actions of one trans. followed by all the actions of another transaction and so on.

Ex.

$T_1$	$T_2$
Read(A, t)	
$t += 100$	
Write(A, t)	
Commit	
	$T_1; T_2$
	Read(A, s)
	$s += 1.1$
	Write(A, s)
	Commit

$T_1$	$T_2$
Read (A, t) t += 100 Write (A, t) Commit	Read (A, s) s *= 1.1 Write (A, s) Commit

$T_2; T_1$

Each schedule might have a different impact on DB.

Say  $A_0$  value of A before schedule.

$$T_1; T_2 \Rightarrow A = 1.1 (A_0 + 100)$$

$$T_2; T_1 \Rightarrow A = 100 + 1.1 A_0$$

### Serializable Schedule.

A schedule S is serializable if there exists a serial schedule S' of the same transactions such that for every initial state of the DB, the effect of S and S' is the same.

$T_1$   
Read (A, t)  
  
 $t + 100$   
Write (A, t)

Commit;

$T_2$   
  
Read (A, s)  
  
 $s * = 1.1$   
Write (A, s)  
Commit

Effect of schedule:

$A = 1.1 A \neq \text{effect of } T_1; T_2 \text{ or } T_2; T_1.$

$\Rightarrow$  non-serializable.



Another schedule:

$T_1$   
Read (A, t)  
 $t += 100$   
Write (A, t)

Commit;

$T_2$   
  
Read (A, s)  
 $s *= 1.1$   
Write (A, s)  
Commit

Serializable:

Equivalent to  $T_1; T_2$

To model transactions we only care about  
Read, Write, Commit, Rollback.

We can rewrite the schedule above as:

$R_1(A), W_1(A), R_2(A), W_2(A), C_2, C_1$   
Use  $A_i$  for rollback (abort).

The job of the DBMS is to only  
allow serializable schedules.

## Anomalies due to interleaved execution

There are 3 main ways in which 2 interleaved transactions can leave the DB in an inconsistent state.

Two actions on the same data object conflict if at least one of them is a write.

### 1) Read Uncommitted Data

#### Dirty Read

- $T_1$  writes to object A
- Before  $T_1$  ends  $T_2$  reads A.

### 2) Unrepeatable Read

- $T_1$  reads an object
- $T_2$  changes the value of the same object before  $T_1$  ends.

### 3) Blind Write (Overwriting not committed data)

- $T_1$  updates the value of A
- $T_2$  overwrites the value of A without   
 before reading it, before  $T_1$  ends.

⇒ Lost update.

Ex:

$T_1$  moves 100 from A to B

$T_2$  increases both A and B by 10%

Serial schedules:

$$T_1; T_2 \quad A = 1.1 (A_0 - 100)$$

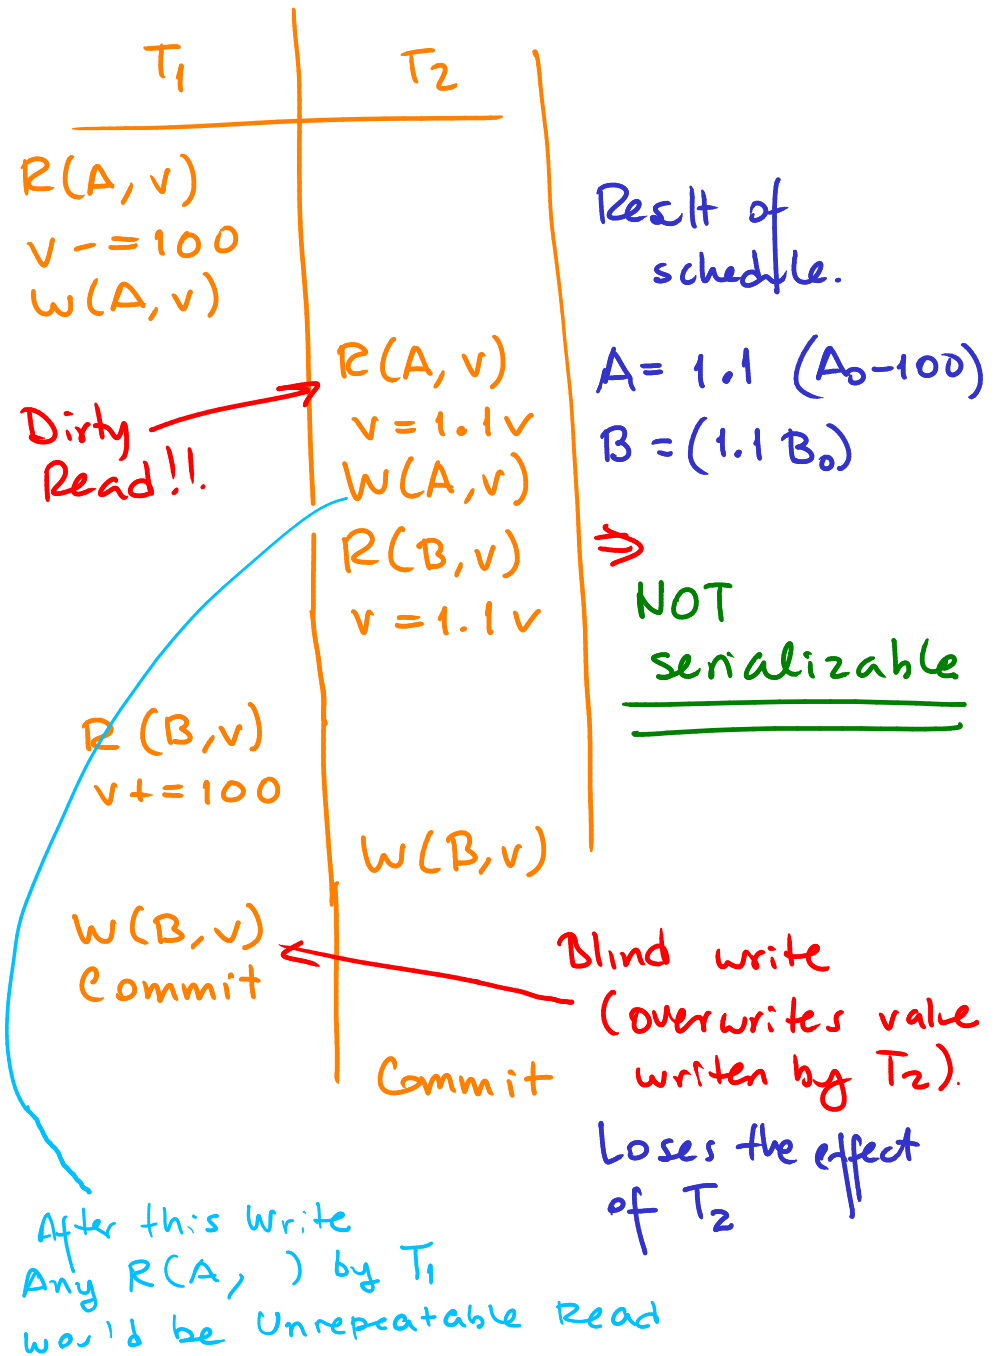
$$B = 1.1 (B_0 + 100)$$

or

$$T_2; T_1 \quad A = (1.1 A_0) - 100$$

$$B = (1.1 B_0) + 100$$

Ex: A non serial schedule:



A transaction with anomalies might  
still result in a serializable schedule.

but...

the DBMS will not create transactions  
with anomalies...unless indicated  
otherwise. (See next notes)